

Market Intelligence Decision-Making System: Predicting Stock with ML

Rahul Chettri, Karthikeyan Sugavanan, Alexis Musaelyan
Northeastern University

{chettri.ra, ksugavanan, amusaelyan}@northeastern.edu

Abstract—This research presents a machine learning-driven **Market Intelligence Decision-Making System** designed to enhance market entry strategies using **S&P 500 data**. The system integrates **predictive modeling** and **Large Language Models (LLMs)** to analyze historical trends, economic indicators, and market conditions. By combining **quantitative analysis** with **natural language processing**, we aim to transform raw financial data into actionable insights, enabling **data-driven investment decisions** in volatile markets.

Index Terms—Machine Learning, Market Prediction, LSTM, S&P 500, Predictive Analytics, AI for Finance

I. INTRODUCTION

Our approach leverages **Long Short-Term Memory (LSTM)**, **XGBoost**, and **Convolutional Neural Networks (CNNs)** to build a comprehensive stock market prediction system. Each model serves a distinct purpose in addressing different aspects of market forecasting.

II. JUSTIFICATION FOR MODEL SELECTION AND COMPARISON

A. LSTM for Time-Series Dependencies

LSTM is particularly well-suited for capturing **long-term dependencies in sequential data**, making it an ideal choice for **financial time-series forecasting**. Since stock prices often exhibit trends influenced by historical movements, LSTM helps identify **underlying temporal relationships** that traditional models might overlook.

B. XGBoost for Feature-Driven Predictions

XGBoost is a powerful **gradient-boosting algorithm** that excels when the dataset is well-structured and engineered with relevant features. By incorporating **technical indicators** such as **moving averages**, **Relative Strength Index (RSI)**, and **Bollinger Bands**, we aim to enhance predictive performance.

Comparing XGBoost with LSTM allows us to evaluate whether a **feature-rich tabular data-driven model** (XGBoost) can outperform a **pure time-series approach** (LSTM). This comparison helps us identify the most **generalizable and effective model** for future implementations.

C. CNN for Pattern Recognition in Stock Charts

In addition to numerical predictions, we explored **Convolutional Neural Networks (CNNs)** to identify **chart patterns** such as **double tops**, **head and shoulders**, and other common stock market formations.

CNNs excel at recognizing **spatial patterns**, making them a natural fit for detecting **visual trends** in stock price movements. By implementing CNNs, we can assess the **effectiveness of deep learning in technical pattern recognition**, adding another layer of intelligence to our market prediction framework.

III. INSIGHTS FROM MODEL COMPARISON

- By comparing **LSTM and XGBoost**, we can determine whether **sequential memory-based learning** (LSTM) or **feature-engineered tabular learning** (XGBoost) provides more reliable market predictions.
- Evaluating **CNN-based pattern recognition** helps us understand whether **visual stock trends** can be effectively incorporated into a predictive framework.
- This comparison ensures we build a **generalized, robust model** that integrates the **best of time-series forecasting, feature engineering, and deep learning-based pattern recognition**.
- This **multi-model approach** allows us to combine **structured feature-based predictions, sequence learning, and pattern recognition**, making our system more **adaptive and intelligent** in financial market forecasting.

IV. PROBLEM STATEMENT: A MULTI-MODEL APPROACH FOR STOCK MARKET PREDICTION

A. Defining the Problem

Stock market forecasting is inherently complex due to its **dynamic, non-linear, and volatile nature**. Traditional statistical models struggle to capture both **temporal dependencies** and **pattern-based insights** effectively. To address this, we propose a **hybrid predictive framework** leveraging **Long Short-Term Memory (LSTM)**, **Extreme Gradient Boosting (XGBoost)**, and **Convolutional Neural Networks (CNNs)**—each serving a specialized role in market prediction.

B. Machine Learning Category

This study falls under the categories of:

- **Time-series forecasting**: Using past trends to predict future price movements.
- **Regression analysis**: Estimating continuous stock price values.
- **Pattern recognition**: Identifying key stock chart formations.

C. Models Used

Each model in our hybrid approach is tailored for a specific function:

- **LSTM**: Captures **long-term dependencies** in sequential stock price movements.
- **XGBoost**: Learns from **engineered features** (technical indicators) to predict price fluctuations.
- **CNN**: Identifies key **chart patterns** (e.g., head and shoulders, double tops) in stock price visualizations.

V. SIGNIFICANCE

This research bridges the gap between **sequential learning, feature-driven analysis, and technical pattern recognition**, ensuring a **robust, multi-faceted market forecasting model**.

A. Academic Relevance

- Contributes to **financial machine learning** by evaluating how different architectures (**sequence-based, feature-based, and pattern-based**) complement each other.
- Assesses the effectiveness of **LSTM vs. XGBoost vs. CNN** in stock market prediction.

B. Industry Applications

- **Quantitative traders, hedge funds, and financial analysts** can leverage this hybrid model for **risk assessment, trading strategies, and market analysis**.
- Can be integrated into **algorithmic trading systems** for automated decision-making.

C. Real-World Impact

- **Retail investors** benefit from a **more reliable prediction system** that combines both **historical trends and real-time market patterns**.
- Helps reduce **uncertainty in investment decisions**, leading to better portfolio management.

VI. MODEL SELECTION AND JUSTIFICATION

A. LSTM for Time-Series Dependencies

LSTM is particularly well-suited for capturing **long-term dependencies in sequential data**, making it an ideal choice for **financial time-series forecasting**. Since stock prices often exhibit **trends influenced by historical movements**, LSTM helps identify **underlying temporal relationships** that traditional models might overlook.

B. XGBoost for Feature-Driven Predictions

XGBoost is a powerful **gradient-boosting algorithm** that excels when the dataset is well-structured and engineered with relevant features. By incorporating **technical indicators** such as **moving averages, Relative Strength Index (RSI), and Bollinger Bands**, we aim to enhance predictive performance.

Comparing XGBoost with LSTM allows us to evaluate whether a **feature-rich tabular data-driven model** (XGBoost) can outperform a **pure time-series approach** (LSTM). This comparison helps us identify the most **generalizable and effective model** for future implementations.

C. CNN for Pattern Recognition in Stock Charts

In addition to numerical predictions, we explored **Convolutional Neural Networks (CNNs)** to identify **chart patterns** such as **double tops, head and shoulders**, and other common stock market formations.

CNNs excel at recognizing **spatial patterns**, making them a natural fit for detecting **visual trends** in stock price movements. By implementing CNNs, we can assess the **effectiveness of deep learning in technical pattern recognition**, adding another layer of intelligence to our market prediction framework.

VII. INSIGHTS FROM MODEL COMPARISON

- By comparing **LSTM and XGBoost**, we can determine whether **sequential memory-based learning** (LSTM) or **feature-engineered tabular learning** (XGBoost) provides more reliable market predictions.
- Evaluating **CNN-based pattern recognition** helps us understand whether **visual stock trends** can be effectively incorporated into a predictive framework.
- This comparison ensures we build a **generalized, robust model** that integrates the **best of time-series forecasting, feature engineering, and deep learning-based pattern recognition**.
- This **multi-model approach** allows us to combine **structured feature-based predictions, sequence learning, and pattern recognition**, making our system more **adaptive and intelligent** in financial market forecasting.

VIII. DATA COLLECTION AND PREPARATION

A. Data Sources

The stock market data is sourced from **Yahoo Finance** (yfinance API), providing **10 years of historical daily stock prices**. The dataset consists of:

- **Date**
- **Open, High, Low, Close, and Volume**

B. Feature Selection

- **Close Price**: The primary feature for model training, as it reflects the overall market sentiment at the end of each trading day.
- **High and Low Prices**: Excluded due to excessive volatility.
- **Open Price**: Omitted since it is highly influenced by pre-market events and news.

C. Data Description

- **Timeframe**: 10 years of daily stock price data.
- **Key Feature**: Close Price (Used for all models: LSTM, XGBoost, and CNN).
- **Additional Features (for XGBoost & CNN)**:
 - Moving Averages (100-day, 200-day, etc.)
 - Peak-Trough Detection for CNN
 - Momentum-based Indicators (RSI, Bollinger Bands, etc.)

D. Handling Missing Values

Although the dataset itself does not contain missing values, certain preprocessing steps introduce NaNs:

1) Moving Averages & Feature Engineering (XGBoost):

- The calculation of **100-day and 200-day Moving Averages** leads to missing values at the beginning (e.g., first 100 days for MA₁₀₀).
- These NaNs are handled by removing early data points where features are incomplete or by extending the dataset to ensure valid values exist.

2) Correction Strategy:

- Cutting the first 200 rows (if 200-day MA is used) to ensure complete feature values for all data points.
- Alternatively, using a tail-extension approach by replicating early values to maintain consistency.

IX. PREPROCESSING STEPS

A. Time Series Feature Engineering (For XGBoost)

- **Moving Averages (MA₅₀, MA₁₀₀, MA₂₀₀):** Smooth out fluctuations and capture long-term trends.
- **Exponential Moving Average (EMA₅₀, EMA₁₀₀):** Provides greater weight to recent prices.
- **Relative Strength Index (RSI):** Identifies overbought/oversold conditions.
- **Bollinger Bands:** Measures volatility around the moving average.

B. Peak Identification for CNN

Using `scipy.signal.find_peaks()` to detect:

- **Local Maxima (Peaks)** → Represents potential resistance levels.
- **Local Minima (Troughs)** → Represents support levels.

Encoding Peaks & Troughs into CNN Input Data:

- Converted into binary peak signals embedded into stock trend images.
- CNN is trained to recognize historical peak-based patterns for trend forecasting.

C. Time-Series Data Preparation for LSTM

- Convert Close Prices into **Windowed Sequences** (e.g., past 50 days used to predict the next day's price).
- Normalize Data (**MinMaxScaler**) for better convergence in neural networks.
- **Train-Test Split:**
 - 80% training, 20% testing while maintaining chronological order.

X. WHY THIS APPROACH?

- **XGBoost:** Leverages historical trends and moving averages to predict future movements.
- **LSTM:** Captures sequential dependencies in stock price changes.
- **CNN:** Detects chart patterns using peaks and troughs, allowing the model to recognize technical formations like head & shoulders, double tops, and support/resistance zones.

XI. SELECTION OF MACHINE LEARNING MODELS

A. Model Consideration

For this project, we explored various machine learning and deep learning models to predict stock prices and identify market trends. The selection process involved comparing **traditional machine learning techniques** and **advanced deep learning models** to determine the most effective approach.

B. Models Considered and Justification

1) Long Short-Term Memory (LSTM) – Time-Series Dependency Modeling: **Reason for Consideration:**

- LSTM is designed for sequential data and is well-suited for stock price forecasting.
- It captures long-term dependencies and trend patterns in stock movements.

Advantages:

- Learns patterns in historical price trends.
- Handles vanishing gradient issues better than standard RNNs.

Challenges:

- Computationally expensive compared to traditional models.
- Performance depends heavily on sequence length and hyperparameter tuning.
- Prone to overfitting if not properly regularized.

2) Extreme Gradient Boosting (XGBoost) – Feature-Based Learning: **Reason for Consideration:**

- XGBoost is highly effective for structured tabular data, allowing the integration of engineered features such as moving averages, RSI, and Bollinger Bands.

Advantages:

- Fast training and efficient for numerical data.
- Handles missing values and outliers well.
- Provides feature importance insights for model interpretability.

Challenges:

- Does not inherently consider temporal dependencies.
- Requires extensive feature engineering.

3) Convolutional Neural Networks (CNN) – Pattern Recognition in Stock Trends: **Reason for Consideration:**

- CNNs are effective for identifying visual patterns in stock price trends, such as peaks, troughs, and formations like head & shoulders and double tops.

Advantages:

- Recognizes technical chart patterns.
- Generalizes well to unseen stock data when trained properly.
- Less sensitive to short-term fluctuations compared to LSTM.

Challenges:

- Requires transforming numerical data into image-like representations.
- Longer training time compared to XGBoost.

C. Final Model Selection: CNN

Reason for Selection:

- CNN performed well in recognizing stock patterns.
- It captured technical formations effectively, reducing noise from short-term price fluctuations.
- Provided strong interpretability by visualizing market trend patterns.

XII. EVALUATION METRICS USED FOR COMPARISON

To assess the performance of the models, the following evaluation metrics were considered:

- **Mean Absolute Error (MAE):** Measures the average absolute differences between predicted and actual values.
- **Root Mean Squared Error (RMSE):** Penalizes larger errors more significantly than MAE.
- **Mean Squared Error (MSE):** Provides an overall measure of prediction deviation.
- **R² Score:** Measures how well predictions fit actual data.
- **Accuracy (%):** Used primarily for CNN-based classification of stock trend patterns.

XIII. MODEL DEVELOPMENT AND TRAINING

A. Architecture and Configuration

1) LSTM:

- **Number of Features:** 1 (Close price)
- **Data Split:** 70% Training, 30% Testing

2) XGBoost:

- **Number of Features:** 6 (Technical indicators such as moving averages, RSI, Bollinger Bands)
- **Data Split:** 80% Training, 20% Testing

3) CNN:

- **Number of Features:** Close price with detected peaks and dips
- **Data Split:** 80% Training, 20% Testing

B. Training Process

- **Early Stopping:** Implemented for LSTM and CNN to stop training when validation loss no longer improves.
- **Dropout Layers:** Added to LSTM and CNN architectures to prevent overfitting.
- **Batch Normalization:** Used in CNN to stabilize training.
- **Cross-Validation:** Applied in XGBoost to optimize performance and reduce variance.

C. Hyperparameter Tuning

- **Grid Search:** Used for all models to find the optimal batch size, learning rate, and number of layers.
- **Dropout Rate:** Tuned to reduce overfitting.
- **Epoch Selection:** Adjusted based on validation loss to prevent unnecessary overtraining.

XIV. EVALUATION AND COMPARISON

A. Key Performance Metrics

To evaluate and compare the models, the following performance metrics were used:

- **Root Mean Squared Error (RMSE):** Measures the prediction error, giving higher weight to large deviations.
- **Mean Squared Error (MSE):** The average squared difference between predicted and actual values.
- **Mean Absolute Error (MAE):** Measures the absolute difference between predicted and actual values.
- **R² Score:** Measures how well the model explains the variance in stock prices, with values closer to 1 indicating better performance.
- **Accuracy (%):** Used for CNN to evaluate its effectiveness in detecting stock market patterns.

B. Model Performance Comparison

1) CNN Model Performance:

- **RMSE:** 357.96
- **Accuracy:** 97.42%
- **Specialization:** CNN was primarily used for pattern recognition, identifying key stock trading patterns such as *head & shoulders*, *double tops*, and *trend reversals*.

2) **LSTM Model Performance:** The best performance was achieved at 100 epochs with batch size 32:

- **RMSE:** 99.10
- **MSE:** 9,044.55
- **R² Score:** 0.9483
- **MAE:** 90.54

LSTM significantly improved with more training, reducing RMSE from **242.21 (30 epochs)** to **99.10 (100 epochs)**. It effectively captured **sequential dependencies** in stock price movements.

3) **XGBoost Model Performance:** XGBoost used engineered features (e.g., moving averages, RSI, Bollinger Bands) and had consistent improvements across epochs. The best results at 100 epochs with batch size 32:

- **RMSE:** 105.00
- **MSE:** 11,025.00
- **R² Score:** 0.9700
- **MAE:** 85.00

XGBoost performed slightly worse than LSTM in RMSE but had a **higher R² score**, indicating better variance explanation. It was computationally efficient and handled **feature-based learning** effectively.

C. Final Model Selection: CNN

Although LSTM and XGBoost were strong candidates for numerical forecasting, CNN was chosen due to its superior **accuracy (97.42%)** in stock pattern recognition.

Reasons for Selection:

- More robust against **short-term volatility** than LSTM.
- Effectively recognized **technical patterns**, aiding traders in making strategic decisions.

- Superior at detecting **trend reversals and high-probability trading setups**.
- Pattern-based trading signals can complement numerical forecasting approaches.